

Databases and Information Systems

CS303

Recovery Management
05-10-2023

Potential Failures

- Disk crash
 - Power outage
 - Software error
 - Natural disaster
 - Security breach
-
- How to ensure atomicity and durability during system failures?

Failure Classification

- Transaction failure
 - Logical error : No loss of data
 - System error : Transaction can be re-executed
- System crash : Due to bug in database software or OS
 - Volatile contents are erased but non-volatile contents remain intact (fail-stop assumption)
- Disk failures :
 - Backups are used to restore the data

Recovery Algorithm

- Two functions:
 - Actions taken during normal transaction processing to ensure that enough information exists to allow recovery from failures.
 - Actions taken after a failure to recover the database contents to a state that ensures database consistency, transaction atomicity, and durability.

Stable Storage Implementation

- Recovery System is as good as its implementation of **stable storage**
 - Replicate the needed information in several disk with independent failure modes
 - Update the information in a controlled manner to ensure that failure during data transfer does not damage the needed information.
- Simplest way to do this is to have **mirrored disks**
 - Two disks that mirror each other in the same drive
 - More copies implies less performance, but greater reliability
 - Does not protect against natural disasters
- Use **remote backups** stored in a different location
 - Transfer data through network (As soon as an output operation is complete)

Stable-storage Implementation

- System should detect data-transfer failure and invoke recovery procedure to restore the block to consistent state
 - To do this system maintains 2 logical database blocks
(in same location for mirrored disks, one local another remote for remote backups)
 - Write information onto the first physical block
 - Write information onto the second physical block
 - Success only after second write is completed
- If failure happens in the middle, two copies are inconsistent with each other.

Stable-storage Implementation

- In recovery phase:
 - If both copies contain no detectable errors (using checksum) then we are good
 - If one has error then its data is replaced by the other copy
 - If both contain no detectable error but contents are different then first copy's data is written to second copy
- Ensures that writes to stable storage succeeds completely or there is no change at all
- Comparing each block for detectable errors is costly
 - Maintain small nonvolatile RAM that remembers which block writes are in progress.
 - Check only these blocks
- Same strategy can be used to have multiple copies of stable storage

Data Access

- **Input / Output** requires bring data from non-volatile memory to volatile memory
- Transferred in blocks (that contains many data items)
 - Assume no overspill
- **Physical blocks:** Blocks residing in disk
- **Buffer blocks :** Blocks residing in buffer
- **Disk buffer :** Area of memory where blocks are stored temporarily
- Block movement involves:
 - **input(B)** transfers the physical block B to main memory.
 - **output(B)** transfers the buffer block B to the disk, and replaces the appropriate physical block

Data Access

- Each transaction T_i has a private work area in which copies of data items accessed and updated by T_i are kept.
 - The system creates this work area when the transaction is initiated
 - Removes it when the transaction either commits or aborts
 - Each data item X kept in the work area of transaction T_i is denoted by x_i
 - Transaction T_i interacts with the database system by transferring data to and from its work area to the buffer.
- We transfer data by these two operations:
 - **read(X)** assigns the value of data item X to the local variable x_i .
 - If block B_x on which X resides is not in main memory, it issues $\text{input}(B_x)$
 - It assigns to x_i the value of X from the buffer block.
 - **write(X)** assigns the value of local variable x_i to data item X in the buffer block.
 - If block B_x on which X resides is not in main memory, it issues $\text{input}(B_x)$.
 - Assigns the value of x_i to X in buffer B_x
- Some data of B_x might still be accessed, so **X is not written to disk immediately**
- **If system crashed before $\text{output}(B_x)$ is executed then X is never written to main memory**

Recovery and Atomicity

- Suppose transaction T_i transfers 50 rupees from A to B (with $A = 1000$, $B = 2000$ initially).
- If system fails **after output(A) but before output(B)**
 - When system restarts, $A = 950$ and $B = 2000$
 - **Impossible** to know which outputs were executed and which were not, just by looking at the database state
- If T_i performs several outputs, failure may occur when some (but not all) updates have been transferred to the disk
- **Atomicity** : Should keep track of changes and either complete it or rollback
 - Use log records

Log File

- Sequence of log records
- Typical Log record has:
 - Transaction Identifier
 - Data-item identifier (location of the data in the disk)
 - Old value
 - New value
 - $\langle T_j X_j V_1 V_2 \rangle$
 - $\langle T_i \text{start} \rangle$
 - $\langle T_i \text{commit} \rangle$
 - $\langle T_i \text{abort} \rangle$
- Log must be stored in stable storage to help disk recovery
- Log may become huge - How to safely erase log ?

Database Modification

- How does a **transaction modify the data?**
 - The transaction performs some computations in its own private part of main memory
 - The transaction modifies the data block in the disk buffer in main memory holding the data item
 - The database system executes the output operation that writes the data block to disk
- **Differed modification** : Transaction does not modify the database until it is committed
 - Needs to keep track of all updates locally
- **Immediate modification** : Transaction modifies the database while it is still active

Database Modification

- Recovery Algorithm should consider the following factors :
 - The possibility that a transaction may have committed although some of its database modifications exist only in the disk buffer in main memory and not in the database on disk.
 - The possibility that a transaction may have modified the database while in the active state and, as a result of a subsequent failure, may need to abort.
- Operations:
 - Undo : Using log records, set data item to old value
 - Redo : Using log records, set data item to new value

Transaction Commit

- Transaction has **committed** when its commit log record has been output to stable storage
- All earlier log records have already been output to stable storage
- Enough information in the log to ensure that even if there is a system crash, the updates of the transaction can be redone
- If a system crash occurs before a log record $\langle T \text{ commit} \rangle$ is output to stable storage, transaction T will be rolled back.
- The output of the block containing the commit log record is the **only action** that results in a transaction getting committed

Using Log to Undo/Redo

- First consider rollback of normal transactions (no system crash)

```
T0: read(A);  
    A := A - 50;  
    write(A);  
    read(B);  
    B := B + 50;  
    write(B).
```

```
T1: read(C);  
    C := C - 100;  
    write(C).
```

```
<T0 start>  
<T0, A, 1000, 950>  
<T0, B, 2000, 2050>  
<T0 commit>  
<T1 start>  
<T1, C, 700, 600>  
<T1 commit>
```

Using Log to Undo/Redo

- **redo(T)** sets the value of all data items updated by transaction T to the new values.
 - Order of redo matters

Log	Database
$\langle T_0 \text{ start} \rangle$	
$\langle T_0, A, 1000, 950 \rangle$	
$\langle T_0, B, 2000, 2050 \rangle$	
	$A = 950$
	$B = 2050$
$\langle T_0 \text{ commit} \rangle$	
$\langle T_1 \text{ start} \rangle$	
$\langle T_1, C, 700, 600 \rangle$	
	$C = 600$
$\langle T_1 \text{ commit} \rangle$	

Using Log to Undo/Redo

- **undo(T)** restores the value of all data items updated by transaction T to the old values.
 - Also written to log (old values are not important)
 - After undo is completed, **<T abort>** is written to log

Log	Database
<T ₀ start>	
<T ₀ , A, 1000, 950>	
<T ₀ , B, 2000, 2050>	
	A = 950
	B = 2050
<T ₀ commit>	
<T ₁ start>	
<T ₁ , C, 700, 600>	
	C = 600
<T ₁ commit>	

Using Log to Undo/Redo

- If there is a system crash:
 - Transaction T needs to be undone if the log contains the record `<T start>`, but does not contain either the record `<T commit>` or the record `<T abort>`
 - If `<T commit>` is present then redo the operations of T
 - If `<T abort>` is present
 - Then also we should redo the operations of T
 - Because log already has the 'undo' records which will be repeated

Using Log to Undo/Redo

- If there is a system crash:
 - Transaction T needs to be undone if the log contains the record `<T start>`, but does not contain either the record `<T commit>` or the record `<T abort>`
 - If `<T commit>` is present then redo the operations of T
 - If `<T abort>` is present
 - Then also we should redo the operations of T
 - Because log already has the 'undo' records which will be repeated

- Example : Crash before WRITE(B)

Undo T_0

```
T0: read(A);  
    A := A - 50;  
    write(A);  
    read(B);  
    B := B + 50;  
    write(B).
```

```
<T0 start>  
<T0, A, 1000, 950>  
<T0, B, 2000, 2050>
```

```
T1: read(C);  
    C := C - 100;  
    write(C).
```

Using Log to Undo/Redo

- If there is a system crash:
 - Transaction T needs to be undone if the log contains the record **<T start>**, but does not contain either the record **<T commit>** or the record **<T abort>**
 - If **<T commit>** is present then redo the operations of T
 - If **<T abort>** is present
 - Then also we should redo the operations of T
 - Because log already has the 'undo' records which will be repeated
- Example : Crash before WRITE(C) redo T_0 undo T_1

T_0 : read(A);
 $A := A - 50$;
 write(A);
 read(B);
 $B := B + 50$;
 write(B).

T_1 : read(C);
 $C := C - 100$;
 write(C).

< T_0 start>
< T_0 , A, 1000, 950>
< T_0 , B, 2000, 2050>
< T_0 commit>
< T_1 start>
< T_1 , C, 700, 600>

Using Log to Undo/Redo

- If there is a system crash:
 - Transaction T needs to be undone if the log contains the record `<T start>`, but does not contain either the record `<T commit>` or the record `<T abort>`
 - If `<T commit>` is present then redo the operations of T
 - If `<T abort>` is present
 - Then also we should redo the operations of T
 - Because log already has the 'undo' records which will be repeated
- Example : Crash after `<T1 commit>` redo T₀ redo T₁

T₀: read(A);
A := A - 50;
write(A);
read(B);
B := B + 50;
write(B).

T₁: read(C);
C := C - 100;
write(C).

`<T0 start>`
`<T0, A, 1000, 950>`
`<T0, B, 2000, 2050>`
`<T0 commit>`
`<T1 start>`
`<T1, C, 700, 600>`
`<T1 commit>`

Checkpoints in logs

- Log files are **huge**
 - Redoing all terminated transitions is **safe but time consuming**
- Create **checkpoint operating**
 - **Do not perform any updates** while checkpoint operation is being performed
 - **Output all modified buffer blocks to disk** when checkpoint is performed
- **Checkpoint operation**
 - Output onto stable storage all log records currently residing in main memory.
 - Output to the disk all modified buffer blocks.
 - Output onto stable storage a log record of the form **<checkpoint L>**
 - **L** is a list of transactions active at the time of the checkpoint.

Checkpoints in logs

- < checkpoint L >
 - If a transaction T has terminated before the checkpoint then all database modifications by T are written to database prior to or during the checkpoint operation
 - Such transactions do not need redo
- System can go to the last < checkpoint L > in log file (by searching backwards) and only undo/redo the transactions in L
- System can erase contents of the log before the checkpoint to reclaim space (except the transactions active during the checkpoint)
- Requirement that transactions should not update during checkpoints is bad
 - Use Fuzzy checkpoints