

# Databases and Information Systems

## CS303

---

Indexing and Hashing  
12-10-2023

# B<sup>+</sup> Trees

- Disadvantage of index sets : Performance degrades as file size grows
- B<sup>+</sup> - tree index structure is the most widely used index structures to maintain efficiency despite insertion and deletion of data.
- B<sup>+</sup> - tree index takes the form of a balanced tree in which every path from the root of the tree to a leaf of the tree is of the same length.
  - Each non-leaf node in the tree has between  $n/2$  and  $n$  children, where  $n$  is fixed for a particular tree.

# B<sup>+</sup> Trees : Structure

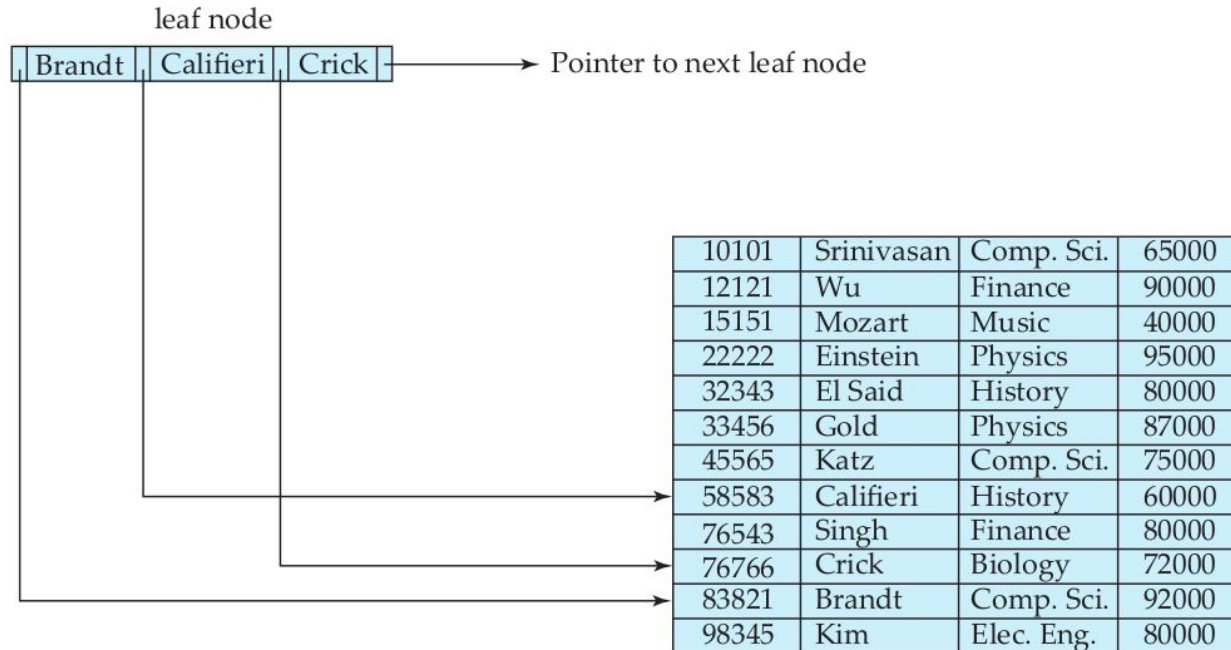
- Let  $n$  be the maximum number of children a node can have
- Each node contains upto  $n-1$  search-keys ( $K_1 \dots K_n$ ) and  $n$  Pointers ( $P_1 \dots P_n$ )
- Search keys in a node are stored in an order
  - If  $i < j$  then  $K_i < K_j$



# B<sup>+</sup> Trees : Structure

- Leaf Node :

- For  $i = 1$  to  $n-1$  Each pointer  $P_i$  in leaf node points to a record in the database
- Pointer  $P_n$  points to the next leaf node



# B<sup>+</sup> Trees : Structure

- Leaf Node :

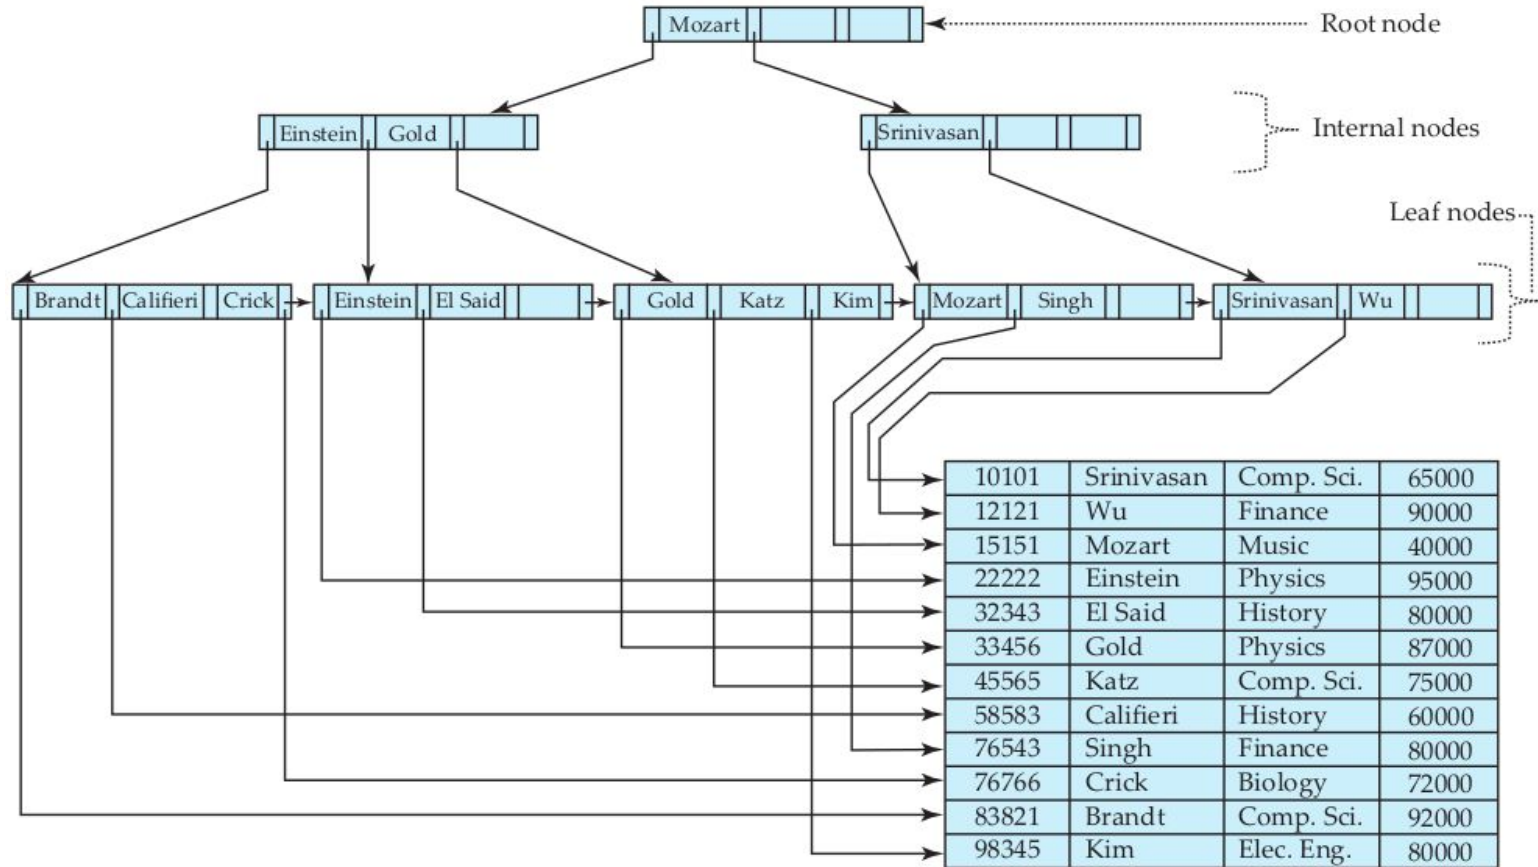
- For  $i = 1$  to  $n-1$  Each pointer  $P_i$  in leaf node points to a record in the database
- Pointer  $P_n$  points to the next leaf node
- Leaf node can have at most  $n$  search-keys and at least  $\lceil (n-1)/2 \rceil$  search-keys
  - With  $k$  search-keys, the node has  $k+1$  pointers
- Values in leaf nodes do not overlap
  - Except if there are duplicates search-key values
- B<sup>+</sup> - tree is a dense index (every search-key appears in some leaf node)

# B<sup>+</sup> Trees : Structure

- Non-Leaf Node :

- Forms multi-level index on the leaf nodes (sparse)
- Can have at most  $n$  search-keys and at least  $\lceil n/2 \rceil$  search-keys
  - With  $k$  search-keys, the node has  $k+1$  pointers
- Root node contains at most  $\lceil n/2 \rceil$  search-keys and at least 2 search-keys
- Every pointer points to a node in the tree

# B<sup>+</sup> Trees : Structure



# Queries on B<sup>+</sup> Trees

**function** find(*value V*)

*/\* Returns leaf node C and index i such that C.P<sub>i</sub> points to first record*

*\* with search key value V \*/*

Set C = root node

**while** (C is not a leaf node) **begin**

Let  $i$  = smallest number such that  $V \leq C.K_i$

**if** there is no such number  $i$  **then begin**

Let  $P_m$  = last non-null pointer in the node

Set C = C.P<sub>m</sub>

**end**

**else if** ( $V = C.K_i$ )

**then** Set C = C.P<sub>i+1</sub>

**else** C = C.P<sub>i</sub> */\* V < C.K<sub>i</sub> \*/*

**end**

*/\* C is a leaf node \*/*

Let  $i$  be the least value such that  $K_i = V$

**if** there is such a value  $i$

**then** return (C,  $i$ )

**else** return null ; */\* No record with key value V exists \*/*

**procedure** printAll(*value V*)

*/\* prints all records with search key value V \*/*

Set done = false;

Set (L,  $i$ ) = find(V);

**if** ((L,  $i$ ) is null) **return**

**repeat**

**repeat**

Print record pointed to by L.P<sub>i</sub>

Set  $i = i + 1$

**until** ( $i >$  number of keys in L **or**  $L.K_i > V$ )

**if** ( $i >$  number of keys in L)

**then** L = L.P<sub>n</sub>

**else** Set done = true;

**until** (done **or** L is null)



# Queries on B<sup>+</sup> Trees

- Can also be used for range queries
- Typically a node has same size of a block (usually 4KB)
- If search key is 12Bytes and Node pointers are 8 Bytes then  $n \sim 200$
- Even with  $n = 100$  and 1 million search keys, we need to access about 4 blocks
- Root node is always accessed (so stays in buffer)
  - So only 3 or fewer blocks are read from the disk

# Difference between B<sup>+</sup> trees and Binary Trees

- Binary trees are stored in memory
- Node size differs by orders of magnitude
- B<sup>+</sup> trees have large branching with short paths
  - Binary trees have small branching with long paths
- Search in Binary tree takes  $O(\log_2 n)$  where  $n$  is the number of nodes
- Search in B<sup>+</sup> trees take  $O(\log_m n)$ 
  - where  $m$  is the max number of children and  $n$  is the number of nodes

# Updating B<sup>+</sup> trees (with n children)

- Insertion

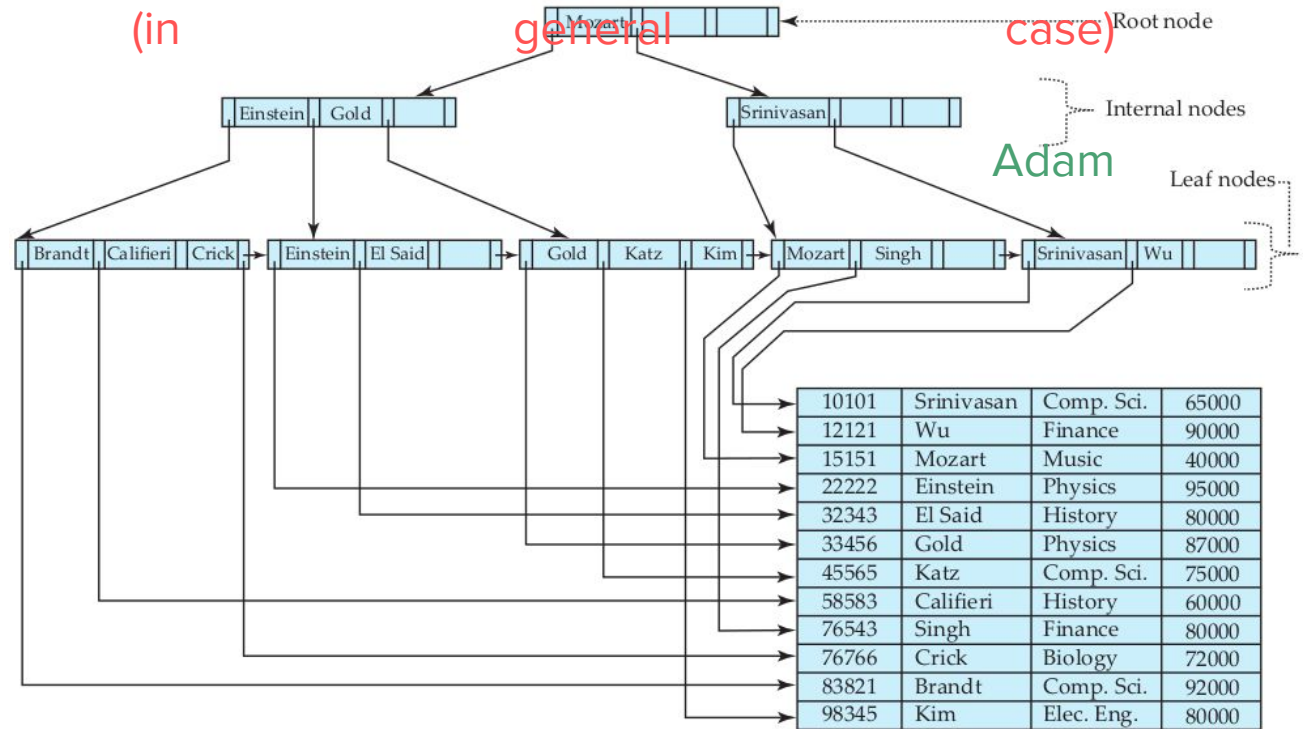
- First find the appropriate place in the leaf where the key should be inserted (using find procedure)
- Suppose the leaf node has fewer than  $n-1$  search keys, then insert the new key in the appropriate position (May have to shift few search-keys and pointers)

- Deletion

- First find the appropriate place in the leaf where the key should be inserted (using find procedure)
- If there are multiple search-keys with same value, go to appropriate record that should be deleted
- Suppose the leaf has more than  $n/2$  search keys, then delete the key and the associated record (may have to shift few search-keys and pointers)

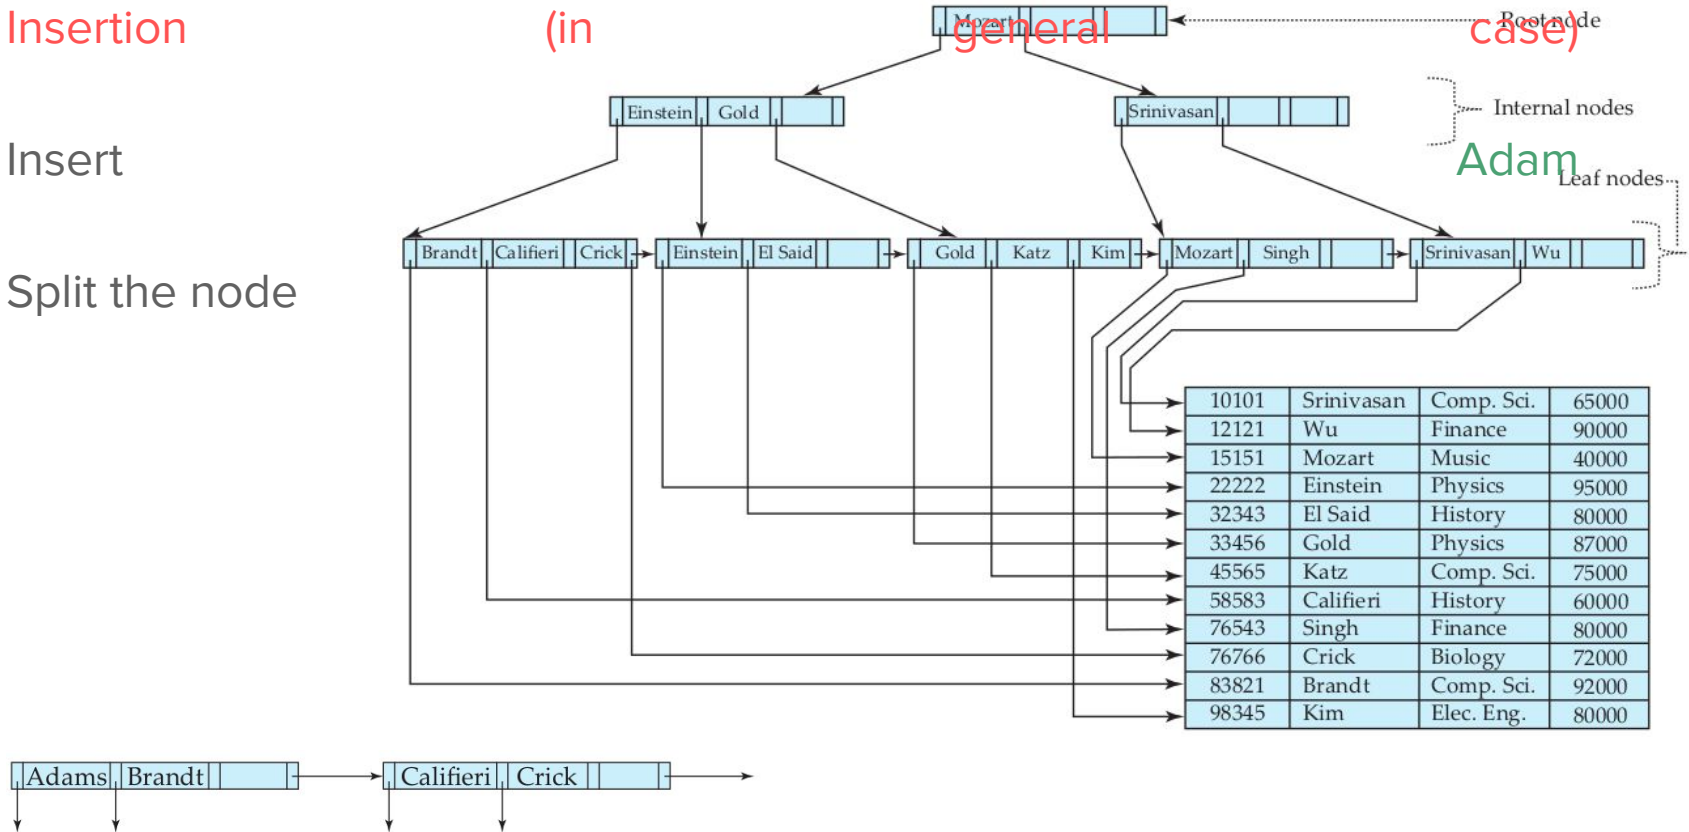
# Updating B<sup>+</sup> trees (with n children)

- Insertion
- Insert
- Split the node



# Updating B<sup>+</sup> trees (with n children)

- Insertion
- Insert
- Split the node


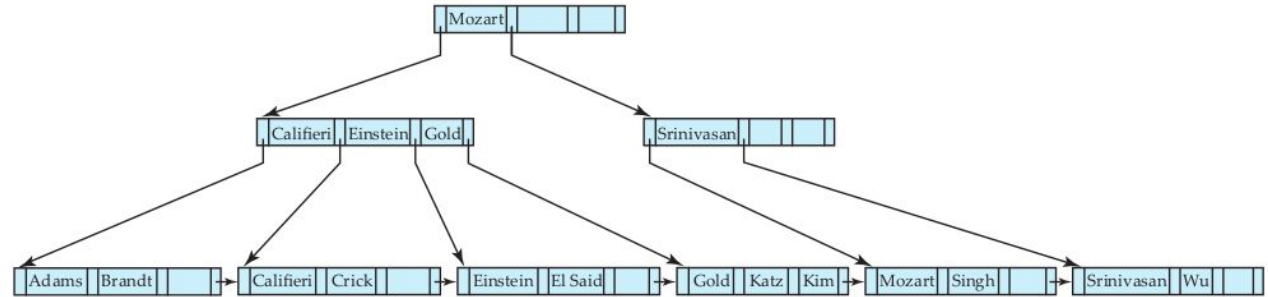


# Updating B<sup>+</sup> trees (with n children)

- Insertion (in general case)

- Insert

- Split the node



# Updating B<sup>+</sup> trees (with n children)

- Insertion

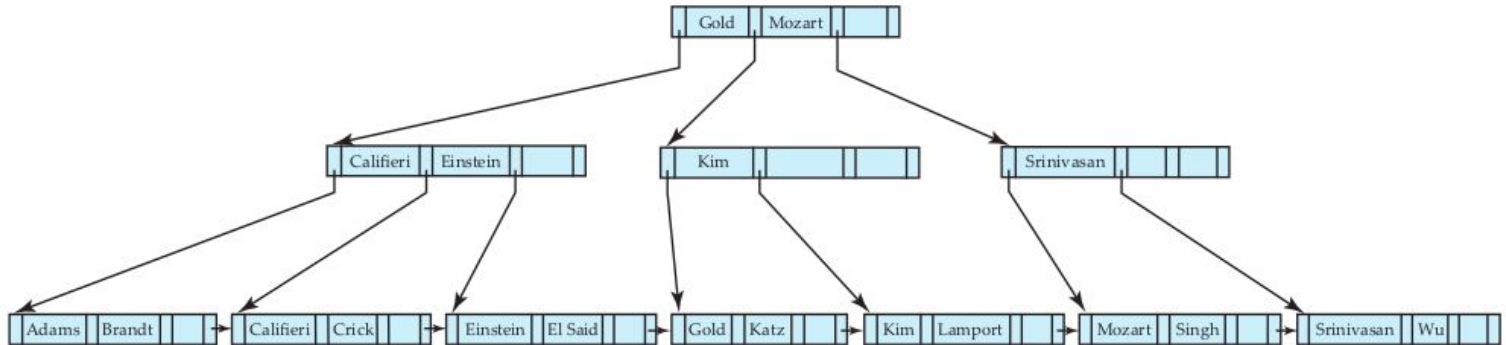
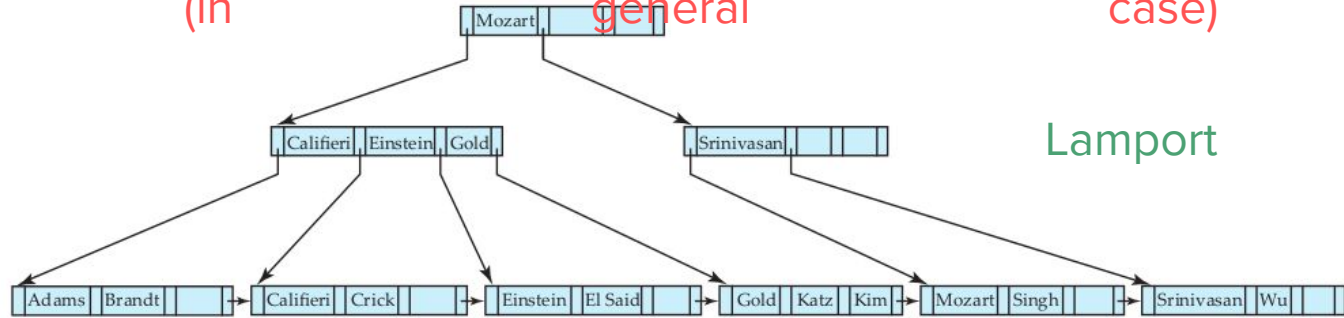
(in

general

case)

- Insert

Lamport



# Insertion in B<sup>+</sup> trees (with n children)

**procedure** insert(*value K, pointer P*)

  if (tree is empty) create an empty leaf node  $L$ , which is also the root

**else** Find the leaf node  $L$  that should contain key value  $K$

  if ( $L$  has less than  $n - 1$  key values)

**then** insert\_in\_leaf ( $L, K, P$ )

**else begin** /\*  $L$  has  $n - 1$  key values already, split it \*/

      Create node  $L'$

      Copy  $L.P_1 \dots L.K_{n-1}$  to a block of memory  $T$  that can  
      hold  $n$  (pointer, key-value) pairs

      insert\_in\_leaf ( $T, K, P$ )

      Set  $L'.P_n = L.P_n$ ; Set  $L.P_n = L'$

      Erase  $L.P_1$  through  $L.K_{n-1}$  from  $L$

      Copy  $T.P_1$  through  $T.K_{\lceil n/2 \rceil}$  from  $T$  into  $L$  starting at  $L.P_1$

      Copy  $T.P_{\lceil n/2 \rceil + 1}$  through  $T.K_n$  from  $T$  into  $L'$  starting at  $L'.P_1$

      Let  $K'$  be the smallest key-value in  $L'$

      insert\_in\_parent( $L, K', L'$ )

**end**

**procedure** insert\_in\_leaf (*node L, value K, pointer P*)

  if ( $K < L.K_1$ )

**then** insert  $P, K$  into  $L$  just before  $L.P_1$

**else begin**

      Let  $K_i$  be the highest value in  $L$  that is less than  $K$

      Insert  $P, K$  into  $L$  just after  $T.K_i$

**end**

**procedure** insert\_in\_parent(*node N, value K', node N'*)

  if ( $N$  is the root of the tree)

**then begin**

      Create a new node  $R$  containing  $N, K', N'$  /\*  $N$  and  $N'$  are pointers

      Make  $R$  the root of the tree

**return**

**end**

  Let  $P = \text{parent}(N)$

  if ( $P$  has less than  $n$  pointers)

**then** insert ( $K', N'$ ) in  $P$  just after  $N$

**else begin** /\* Split  $P$  \*/

      Copy  $P$  to a block of memory  $T$  that can hold  $P$  and ( $K', N'$ )

      Insert ( $K', N'$ ) into  $T$  just after  $N$

      Erase all entries from  $P$ ; Create node  $P'$

      Copy  $T.P_1 \dots T.P_{\lceil n/2 \rceil}$  into  $P$

      Let  $K'' = T.K_{\lceil n/2 \rceil}$

      Copy  $T.P_{\lceil n/2 \rceil + 1} \dots T.P_{n+1}$  into  $P'$

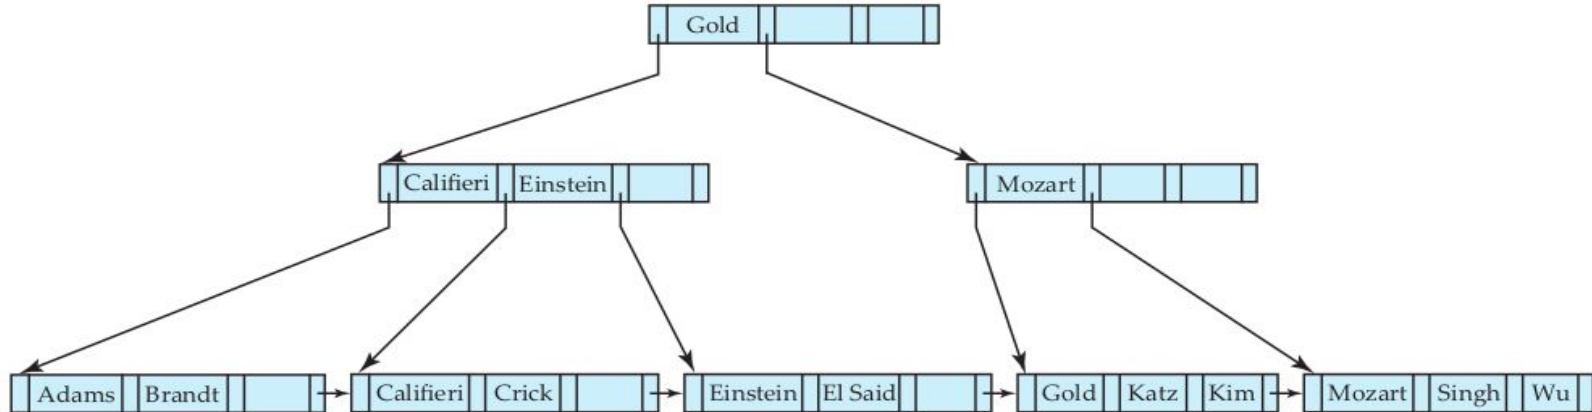
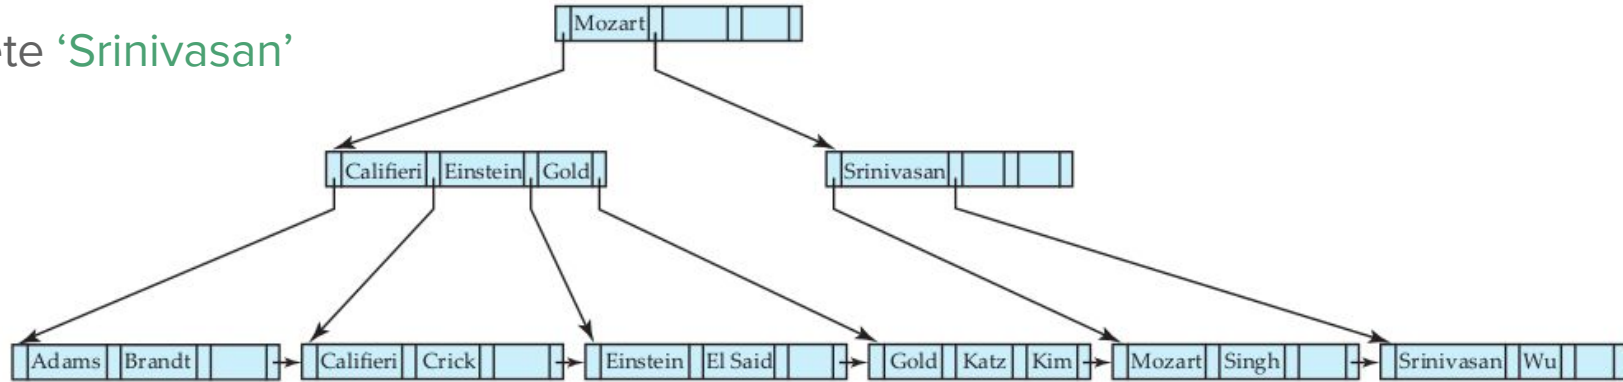
      insert\_in\_parent( $P, K'', P'$ )

**end**




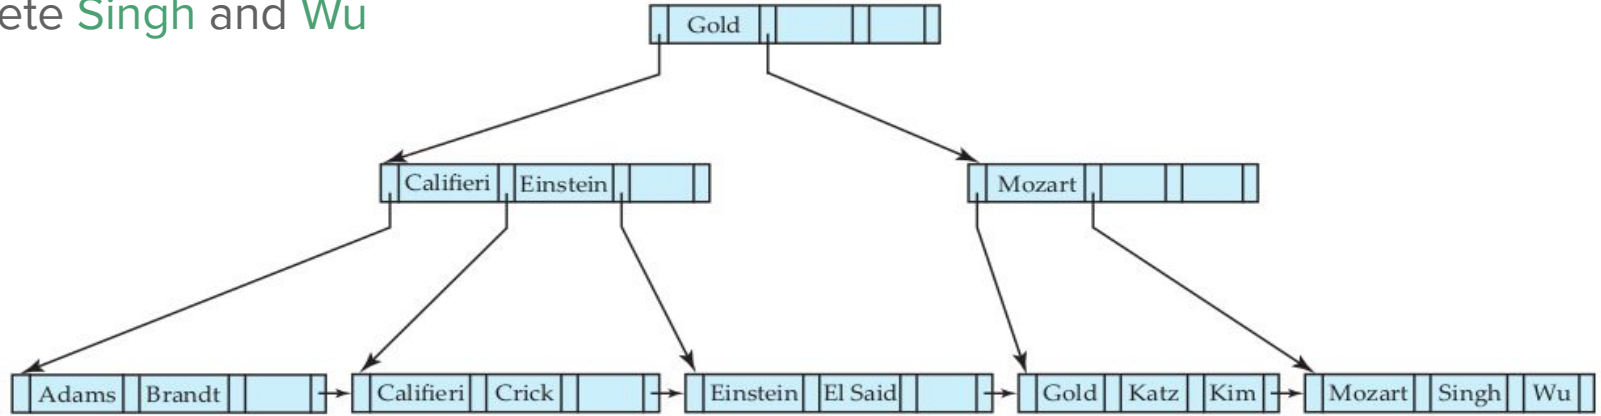
# Deletion in B<sup>+</sup> trees (with n children)

- Delete 'Srinivasan'




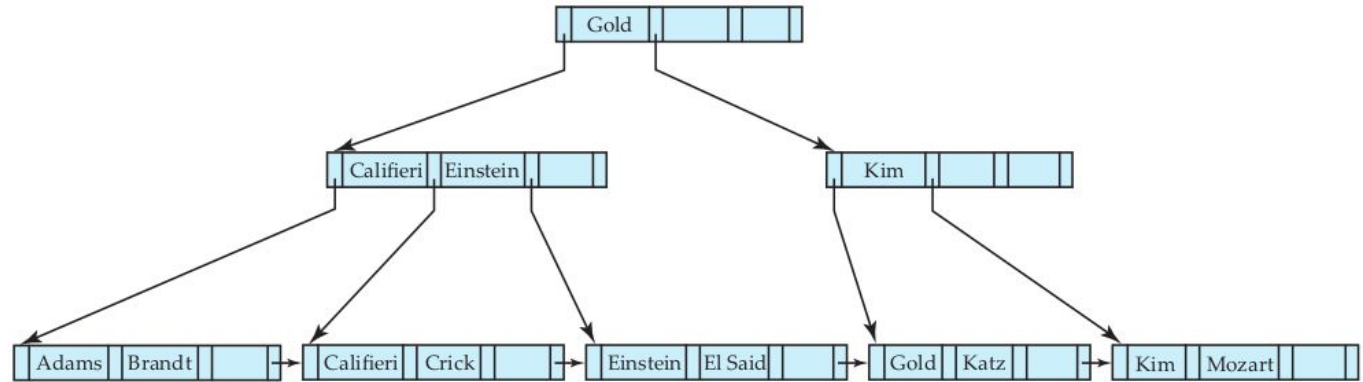
# Deletion in B<sup>+</sup> trees (with n children)

- Delete Singh and Wu



# Deletion in B<sup>+</sup> trees (with n children)

- Delete Gold



# Deletion in B<sup>+</sup> trees (with n children)

```
procedure delete(value K, pointer P)
    find the leaf node L that contains (K, P)
    delete_entry(L, K, P)
```

```
procedure delete_entry(node N, value K, pointer P)
    delete (K, P) from N
    if (N is the root and N has only one remaining child)
    then make the child of N the new root of the tree and delete N
    else if (N has too few values/pointers) then begin
        Let N' be the previous or next child of parent(N)
        Let K' be the value between pointers N and N' in parent(N)
        if (entries in N and N' can fit in a single node)
        then begin /* Coalesce nodes */
            if (N is a predecessor of N') then swap_variables(N, N')
            if (N is not a leaf)
                then append K' and all pointers and values in N to N'
                else append all (Ki, Pi) pairs in N to N'; set N'.Pn = N.Pn
            delete_entry(parent(N), K', N); delete node N
        end
    else begin /* Redistribution: borrow an entry from N' */
        if (N' is a predecessor of N) then begin
            if (N is a nonleaf node) then begin
                let m be such that N'.Pm is the last pointer in N'
                remove (N'.Km-1, N'.Pm) from N'
                insert (N'.Pm, K') as the first pointer and value in N,
                    by shifting other pointers and values right
                replace K' in parent(N) by N'.Km-1
            end
            else begin
                let m be such that (N'.Pm, N'.Km) is the last pointer/value
                    pair in N'
                remove (N'.Pm, N'.Km) from N'
                insert (N'.Pm, N'.Km) as the first pointer and value in N,
                    by shifting other pointers and values right
                replace K' in parent(N) by N'.Km
            end
        end
        else ... symmetric to the then case ...
    end
end
```

# Non unique search-keys

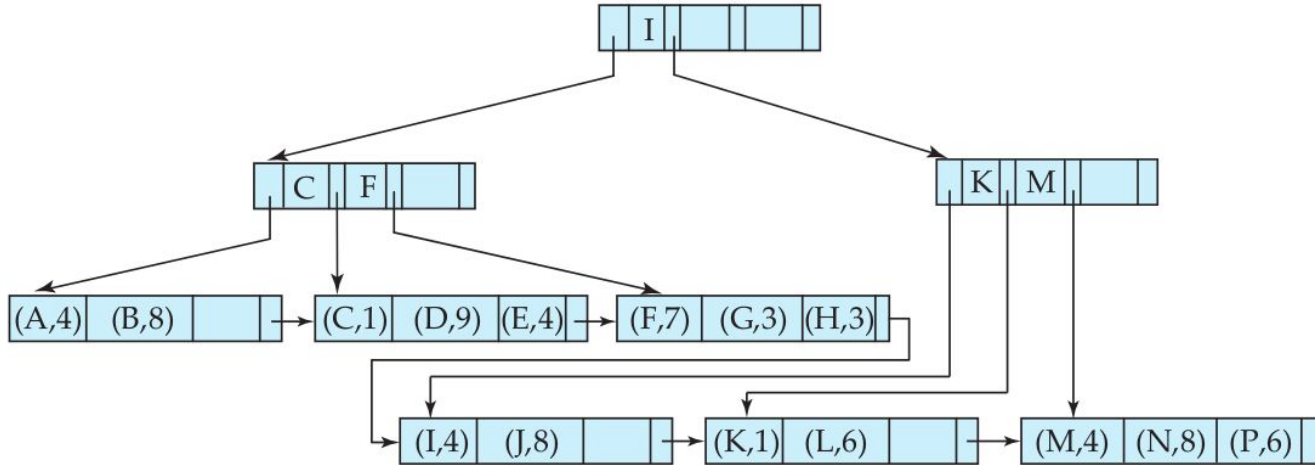
- Keys that are there in more than one tuple of a relation
- If one of them has to be deleted, finding the right node takes time
  - **Solution** : Create composite key with another attribute  
(whose value is unique among all tuples with repeating search-keys)
    - Search query can ignore the new additional attribute
- Alternatively, have unique keys in  $B^+$  tree and leaf points have a bucket
  - Efficient space for the  $B^+$  tree
  - But should deal with variable bucket size
  - More I/O if buckets are stored across blocks
  - Inefficient deletions

# Complexity of B<sup>+</sup> trees Updates

- Insertion and Deletion :  
 $\log_{\lceil n/2 \rceil} (N)$  where  $n$  is the maximum number of children
- In practise, performs fewer I/O than worst case
- For updates, probability of node split is very small (assuming equal search-key distribution)

# B<sup>+</sup> tree Extensions

- B<sup>+</sup> tree File Organization (with respect to clustering search-key):
  - Leaf nodes of the tree stores records instead of pointers
  - Can be used for records with large data types (clob, blob)



# B<sup>+</sup> tree Extensions

- Secondary Indices and File Locations:

- Updates on B<sup>+</sup> tree file organization may change record location even if it is not modified
- All secondary indices to the record should be updated

- Solution:

- Instead of pointers to records, store primary search-key attribute.
- Example: instructor table with ID as clustering search-key index, for department\_name secondary search-key index, for every department\_name store the ID of instructions instead of pointers to records
- No cost for updates, but increases cost of accessing data using secondary index



# B<sup>+</sup> tree Extensions

- **Indexing strings:** If we want to create a B<sup>+</sup> tree index on strings:
  - Strings can be of **variable length**
  - Long string index can lead to small branching (hence long paths in the tree)
- **Branching (fan out) can be increased using prefix compression:**
  - Do not store entire search-key at non-leaf nodes
  - Store only the **prefix of the string** that is sufficient to distinguish between the search-key strings

# B<sup>+</sup> tree Extensions

- Bulk loading of B<sup>+</sup> trees

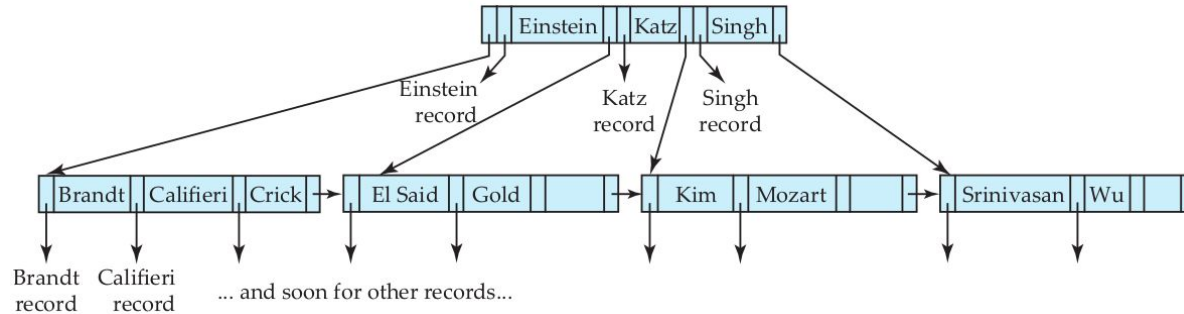
- Creating a secondary index from scratch takes time, since each insertion in the order of clustering index might access different blocks in the created B<sup>+</sup> tree

- Efficient way to do it (Bulk loading)

- Create a temporary relation containing the index entries
- Sort the temporary file using the index entries
- Then all entries that goes into one leaf will appear consecutively
- So all blocks (nodes) need to be brought to main memory once
- With sorted entries, the tree can be built bottom-up.
  - Break the entries into leaf blocks
  - The minimum entry in each block along with the pointer will be the entry for the next level

# B<sup>+</sup> tree Extensions

- B tree index Files
  - Similar to B<sup>+</sup> trees
  - Eliminates redundancy
  - Deletion is more complicated (since deletion can happen in non-leaf nodes too)
  - Most implementations use B<sup>+</sup> trees (but might be called B trees)



# Multiple-Key Access

- Find ID of all instructors in the Finance department whose salary is 80000
  - `SELECT ID FROM instructor WHERE dept_name = 'Finance' AND salary= 80000;`
- Three strategies to use index:
  - Use the index on dept\_name to find all records pertaining to the Finance department. Examine each such record to see whether salary= 80000.
  - Use the index on salary to find all records pertaining to instructors with salary of 80000. Examine each such record to see whether the department name is "Finance".
  - Use the index on dept name to find pointers to all records pertaining to the Finance department. Also, use the index on salary to find pointers to all records pertaining to instructors with a salary of 80,000. Take the intersection of these two sets of pointers.
- Third option seems to be a good option
  - But need not be if the following happens:
    - There are many instructors in Finance department
    - There are many instructors who get 80000 salary
    - There are very few instructors in Finance department who get 80000 salary
- Instead create (dept\_name, salary) together as the search-key index

# Multiple-Key Access

- Suppose we have (dept\_name, salary) as search-key index
- Searching only for 'Finance' department is same as having range query on ('Finance', -  $\infty$ ) to ('Finance', +  $\infty$ )
- There are some drawbacks
  - SELECT ID FROM instructor WHERE dept\_name < 'Finance' AND salary < 80000;
  - Each record is likely to be in a different block because of the ordering of the records
- R-tree is used to handle multiple-key index
  - Generalization of B<sup>+</sup> trees to handle indexing in multiple dimensions

# Multiple-Key Access

- Covering indices

- Stores values of some attributes along with the pointers to the record

- Example: Storing salary value for every ID as search-key

- Same effect as having (ID, salary) as search-key  
but can have larger fanout (hence short trees)