

Databases and Information Systems

CS303

Query Optimization
02-11-2023

Query Optimization

- Goal :
 - Look at all equivalent relational algebraic expressions that are equivalent to the current query
 - Estimate the number of tuples generated at every intermediate step
 - Calculate the cost of evaluation depending on this estimate
 - Choose the expression that has the minimal cost

Recap : Selection size estimation

- $\sigma_{A=a}(r)$

- Estimate of the number of tuples = $n_r / V(A,r)$
 - n_r is the number of tuples in r
 - $V(A,r)$ is the number of distinct values of A that occurs in r
- This assumes all values of A occur with uniform frequency
 - Good approximate in real life
- If histogram data is available then we can get a better estimate

Recap : Selection size estimation

- $\sigma_{A \leq a}(r)$

- Catalog maintains $\min(A,r)$ and $\max(A,r)$ be the lowest and highest values that A can take
- Estimate of the number of tuples
 - If $a < \min(A,r)$ then estimate = 0
 - If $a \geq \max(A,r)$ then estimate = n_r

$$n_r * \frac{a - \min(A,r)}{\max(A,r) - \min(A,r)}$$

- If expression is part of query, value of v may not be available.
 - Then rough estimate is $n_r / 2$

Recap : Selection size estimation

- $\sigma_{\theta_1 \wedge \theta_2 \dots \wedge \theta_n}(r)$

- For each i if the estimate of $\sigma_{\theta_i}(r)$ is s_i then probability of a tuple satisfying θ_i is s_i / n_r
- Estimate

$$n_r * \frac{s_1 * s_2 * \dots * s_n}{n_r^{nr}}$$

- $\sigma_{\theta_1 \vee \theta_2 \dots \vee \theta_n}(r)$

- Estimate

$$n_r * (1 - (1 - s_1/n_r)(1 - s_2/n_r) \dots (1 - s_n/n_r))$$

- $\sigma_{\neg \theta}(r)$ Exercise

Recap : Join size estimation

- Cartesian product $r \bowtie s$ will have $n_r * n_s$ tuples
- Natural Join : Let R and S be the set of attributes of r and s respectively
 - If $R \cap S = \emptyset$ then Natural Join is same as cartesian product
 - If $R \cap S$ is key for R then we can have at most n_s tuples in the natural join
 - If $R \cap S$ is a key neither for R nor for S if $R \cap S = \{ A \}$ then
 - Every tuple t in R produces $n_s / V(A,s)$ many tuples
 - So total estimate is $n_r * n_s / V(A,s)$
 - Reversing the roles of r and s , we get the estimate $n_r * n_s / V(A,r)$
 - Estimates differ if $V(A,r) \neq V(A,s)$ then pick the minimum
- We can also estimate $r \bowtie_{\theta} s$ as $\sigma_{\theta}(r \bowtie s)$

Size estimation for other operations

- Projection : $\Pi_A(r)$ has the estimate $V(A,r)$
 - Since project operation eliminates duplicates
- Set operations :
 - Convert $\sigma_{\theta_1}(r) \cup \sigma_{\theta_2}(r)$ to $\sigma_{\theta_1 \vee \theta_2}(r)$ for estimation
 - If union is over different relations then estimate is the sum of tuples from each operand
 - For intersection of different relations, estimate is the minimum number of both operands
- Outer Join :
 - For left outer join of r and s, estimate is size of $r \bowtie s$ plus size of s
 - Similar for right outer join and full outer join
 - Estimate provides upper bound

Estimation of number of distinct values

- $V(A, \sigma_\theta(r))$
 - Distinct possible values that attribute A can take
 - If θ is of the form $A = a$ then estimate = 1
 - If θ is of the form $A = a_1 \vee A = a_2 \vee \dots \vee A = a_n$ then estimate = n
 - If θ is of the form $A \text{ op } a$ where op is some comparison operator then estimate = $V(A, r) * s/n_r$ where s is the estimate of $\sigma_\theta(r)$
 - In all other cases
 - Assume the selection is independent of A
 - Estimate is the minimum of $V(A, r)$ and estimate of $\sigma_\theta(r)$

Estimation of number of distinct values

- $V(A, r \bowtie s)$
 - If all attributes of A are from R then estimate is minimum of $V(A, r)$ and estimate of $r \bowtie s$
 - If A contains A1 attributes from R and A2 attributes from S then
Estimate is minimum of $V(A1, r) * V(A2-A1, s)$ and $V(A1-A2, r) * V(A2, s)$ and estimate of $r \bowtie s$
- Estimates for aggregates : **sum / count / min / max / average**
 - Exercise

Choice of Evaluation Plan

- **Evaluation plan** defines exactly what algorithm should be used for each operation
 - How the execution of the operations should be coordinated
- **Cost based optimizer** explores all query-evaluation plans that are equivalent to the given query
 - chooses the one with the least estimated cost

Cost based Join Order Selection

- $r_1 \bowtie r_2 \bowtie \dots \bowtie r_n$ can be reordered and the join can be performed
- With $n = 3$ there are 12 different join orderings
- For $n = 10$ we have 17.6 billion possible orderings
- With n different relations to join, how many join orderings are possible?
 - Exercise

| | | | |
|---------------------------------|---------------------------------|---------------------------------|---------------------------------|
| $r_1 \bowtie (r_2 \bowtie r_3)$ | $r_1 \bowtie (r_3 \bowtie r_2)$ | $(r_2 \bowtie r_3) \bowtie r_1$ | $(r_3 \bowtie r_2) \bowtie r_1$ |
| $r_2 \bowtie (r_1 \bowtie r_3)$ | $r_2 \bowtie (r_3 \bowtie r_1)$ | $(r_1 \bowtie r_3) \bowtie r_2$ | $(r_3 \bowtie r_1) \bowtie r_2$ |
| $r_3 \bowtie (r_1 \bowtie r_2)$ | $r_3 \bowtie (r_2 \bowtie r_1)$ | $(r_1 \bowtie r_2) \bowtie r_3$ | $(r_2 \bowtie r_1) \bowtie r_3$ |

Cost based Join Order Selection

- Not necessary to generate all join orderings

- Suppose we want to find the best join ordering of the form

$$(r_1 \bowtie r_2 \bowtie r_3) \bowtie r_4 \bowtie r_5$$

- Choose the best among the 12 orderings of $(r_1 \bowtie r_2 \bowtie r_3)$
- Then join the result with r_4 and r_5
- Total checks = 12 + 12
(instead of 144 with $n = 5$)

- This idea can be used to develop a dynamic-programming algorithm

```
procedure FindBestPlan(S)
  if (bestplan[S].cost  $\neq \infty$ ) /* bestplan[S] already computed */
    return bestplan[S]
  if (S contains only 1 relation)
    set bestplan[S].plan and bestplan[S].cost based on best way of accessing S
  else for each non-empty subset S1 of S such that S1  $\neq$  S
    P1 = FindBestPlan(S1)
    P2 = FindBestPlan(S - S1)
    A = best algorithm for joining results of P1 and P2
    cost = P1.cost + P2.cost + cost of A
    if cost < bestplan[S].cost
      bestplan[S].cost = cost
      bestplan[S].plan = "execute P1.plan; execute P2.plan;
                          join results of P1 and P2 using A"
  return bestplan[S]
```

Cost based optimization of equivalence rules

- Join order optimization technique handles the most common class of queries, which perform an inner join of a set of relations
- But queries use other features which are not addressed by join order selection
 - aggregation, outer join, and nested queries
- Algorithm to generate all possible equivalent relational expressions can be modified to generate all possible query execution plans
 - Example : Join can be annotated as has join, nested-block join etc
- But this is a costly process since it is a Brute-Force algorithm

Cost based optimization of equivalence rules

- Some techniques to make the algorithm efficient:
 - A space-efficient representation of expressions
 - Avoids making multiple copies of the same subexpressions when equivalence rules are applied.
 - Efficient techniques for detecting duplicate derivations of the same expression.
 - A form of dynamic programming based on memoization
 - stores the optimal query evaluation plan for a subexpression when it is optimized for the first time;
 - subsequent requests to optimize the same subexpression are handled by returning the already memoized plan.
 - Techniques to avoid generating all possible equivalent plans
 - By keeping track of the cheapest plan generated for any subexpression up to any point of time
 - Then pruning away any plan that is more expensive than the cheapest plan found so far for that subexpression

Heuristics in Optimization

- A **drawback** of cost-based optimization is the **cost of optimization itself**
- The number of **different evaluation plans for a query can be very large**
 - finding the optimal plan from this set requires a lot of computational effort.
- Hence, **optimizers use heuristics** to reduce the cost of optimization.

Some Heuristics



- Perform selection operation as early as possible
 - $\sigma_{\theta}(r \bowtie s)$ where θ is a condition over r then better to do $\sigma_{\theta}(r) \bowtie s$
- But if r is extremely small compared to s , and if there is an index on the join attributes of s , but no index on the attributes used by θ
 - Then it is probably a bad idea to perform the selection early
 - Performing the selection directly on s require doing a scan of all tuples in s
 - It is probably cheaper, in this case, to compute the join by using the index, and then to reject tuples that fail the selection.

Some Heuristics

- Perform projection operation as early as possible
 - First do selection then do the projection
 - Selection enables the use of index
 - This might also not be the best always
 - Exercise : Come up with a scenario where doing projection later is useful



Some Heuristics

- Some optimizers use particular type of joins called left-deep join orders
- Better for pipelining
- Can find best order (among left-deep joins) faster



Some Heuristics

- Consider **n** different **left-deep join orders**, starting from every relation
- At every step choose the next best relation to perform join



Some Heuristics

- Some optimizers have a budget (time)
- If the budget runs out then return the best query evaluation plan found until then
 - First apply cheap heuristics to get a descent plan
 - In the remaining time apply complicated search to get the best plan

Some Heuristics

- Many applications execute the same query repeatedly but with different values for the constants.
 - Example: a university application may repeatedly execute a query to find the courses for which a student has registered, each time for a different student with a different value for the student ID.
- As a heuristic, many optimizers optimize a query once, with whatever values were provided for the constants when the query was first submitted
 - cache the query plan
- Whenever the query is executed again, the cached query plan is reused
- Optimal plan for the new instance may differ from the optimal plan for the earlier instance
 - but as a heuristic the cached plan is reused
- Caching and reuse of query plans is referred to as Plan caching.

Some Heuristics

- Even with the use of heuristics, cost-based query optimization imposes a substantial overhead on query processing.
- However, the added cost of cost-based query optimization is usually more than offset by the saving at query-execution time.
- The difference in execution time between a good plan and a bad one may be huge, making query optimization essential.
- Most commercial systems include relatively sophisticated optimizers.

Optimizing nested Subqueries (Correlated queries)

- `SELECT name FROM instructor`
`WHERE EXISTS (SELECT * FROM teaches`
`WHERE instructor.ID = teaches.ID AND teaches.year = 2009)`
- **Subquery** can be viewed as a function that takes a parameter (`instructor.ID`) and returns the set of all courses taught in 2009 by instructors with the same ID
- Not optimal to execute the inner query for every tuple from the outer query
- Get better code using **decorrelation**

Optimizing nested Subqueries (Correlated queries)

- `SELECT ... FROM L1
WHERE P1 AND EXISTS (SELECT * FROM L2 WHERE P2)`
- `CREATE TABLE t1 AS SELECT DISTINCT V FROM L2 WHERE P21
SELECT ... FROM L1, t1 WHERE P1 AND P22`
 - V is all attributes used in the condition P₂ from L₂
 - P₂¹ contains predicates in P₂ that talks only about L₂
 - P₂² contains predicates from P₂ that involves the correlation
- `SELECT name FROM instructor
WHERE EXISTS (SELECT * FROM teaches
WHERE instructor.ID = teaches.ID AND teaches.year = 2009)`
- `CREATE TABLE t1 AS SELECT DISTINCT ID FROM teachers WHERE teaches.year = 2009
SELECT name FROM instructor, t1 WHERE t1.ID = instructor.ID`

Optimizing nested Subqueries (Correlated queries)

- Decorrelation is more complicated when
 - The nested subquery uses aggregation
 - The result of the nested subquery is used to test for equality
 - The condition linking the nested subquery to the outer query is not exist
 -
- Decorrelation can be done, but complex
- Optimization of complex nested subqueries is a difficult task
 - It is best to avoid using complex nested subqueries, where possible
 - We cannot be sure that the query optimizer will succeed in converting them to a form that can be evaluated efficiently.

Other optimizations

- Materialized views

- How to update them
 - Incremental view maintenance : Modify only the affected parts of materialized views
 - Immediate view maintenance or Differed view maintenance
- Can be used for decorrelation and optimization

- Top K-optimization (LIMIT K)

- If K is small, no point in computing full result
 - Pipelined plans can be generated in sorted order
 - Estimate what is the highest value on the sorted attributes that will appear in the top-K output
 - Introduce selection predicates that eliminate larger values
 - If extra tuples beyond the top-K are generated they are discarded
 - If too few tuples are generated then the selection condition is changed and the query is re-executed.

Other optimizations

- **Join Minimization** : Dropping a relation from a join without changing the result of the query
 - Can be done if we are creating a view using multiple relations but later using only some of those relations in the query
- Optimization of Updates : **SET ... WHERE**
 - Cannot be done in parallel naively
 - Updates can affect the update query being executed in the presence of index (Halloween problem)
 - Problem can be avoided by executing the queries defining the update first, creating a list of affected tuples, and then updating the tuples and indices
 - Check if the Halloween problem occurs otherwise execute in parallel

Other optimizations

- **Multiquery Optimization and Shared scans** : Happens when a batch of queries are submitted together
 - **Common subexpression elimination** : If same subexpression is used by multiple queries, result can be reused
 - Instead of reading every relation per query, read it once and use it for all relevant queries
- **Parametric Query Optimization (PQO)** :
 - Query is generally optimized for some particular values
(like **SELECT ... FROM... WHERE ID = 1234**)
 - **PQO** optimizes without specifying the value for parameters
 - The optimizer then outputs several plans, each optimal for a different parameter value
 - When a query is submitted with specific values for its parameters
 - Cheapest plan from the set of alternative plans computed earlier is used
 - Finding the cheapest such plan usually takes much less time than reoptimization

Reference:

Database System Concepts by Silberschatz, Korth and Sudarshan
(6th edition)
Chapter 13