

Databases and Information Systems

CS303

Normalization
13-09-2023

Recap : Functional Dependency

- Let R be the set of all attributes of the relational schema $r(R)$ and let $\alpha \subseteq R$ and $\beta \subseteq R$. Given an instance of r , we say $\alpha \rightarrow \beta$ if for any two tuples t_1 and t_2 if $t_1[\alpha] = t_2[\alpha]$ then $t_1[\beta] = t_2[\beta]$
- Example: `inst_dept` (ID, name, salary, dept name, building, budget)
 - `dept_name` \rightarrow building, budget
 - ID, dept_name \rightarrow name, salary, building, budget

Recap : Boyce Codd Normal Form

- A relation schema R is in BCNF with respect to a set F of functional dependencies if, for all functional dependencies in F^+ of the form $\alpha \rightarrow \beta$, where $\alpha \subseteq R$ and $\beta \subseteq R$: at least one of the following holds:
 - $\alpha \rightarrow \beta$ is a trivial functional dependency (that is, $\beta \subseteq \alpha$)
 - α is a superkey for schema R
- A database design is in BCNF if every relational schema of the design is in BCNF

Recap: Boyce Codd Normal Form

General decomposition Rule

- Let $r(R)$ be a schema that is not in BCNF
(There is a non-trivial functional dependency $\alpha \rightarrow \beta$ where α is not a superkey)
- Replace $r(R)$ with two new schemas
 - $\alpha \cup \beta$
 - $(R - (\beta - \alpha))$
- Example: $\text{inst_dept}(\text{ID}, \text{name}, \text{salary}, \text{dept_name}, \text{building}, \text{budget})$
 - $\text{dept_name} \rightarrow \text{building}, \text{budget}$
- Application of the rule might result in smaller relations that are not in BCNF
 - Repeat for procedure for smaller relations till the resulting schema is in BCNF
- Does not preserve functional dependencies

Recap : Third Normal Form

- A relation schema $r(R)$ is in Third Normal Form (3NF) with respect to a set F of functional dependencies if, for all functional dependencies in F^+ of the form $\alpha \rightarrow \beta$, where $\alpha \subseteq R$ and $\beta \subseteq R$: at least one of the following holds:
 - $\alpha \rightarrow \beta$ is a trivial functional dependency (that is, $\beta \subseteq \alpha$).
 - α is a superkey for schema R .
 - Each attribute A in $\beta - \alpha$ is contained in some candidate key of R
- Each such A can be part of different candidate keys
- Any schema that is in BCNF is also in 3NF

Decomposition Algorithms

- To automate the decomposition of the relational schema into normal forms

Recap : Computing F^+

- Use the following Axioms to compute F^+ (Armstrong's axioms)
 - Reflexivity Rule : If $\beta \subseteq \alpha$ then $\alpha \rightarrow \beta$
 - Augmentation Rule : If $\alpha \rightarrow \beta$ and γ is a set of attributes then $\gamma\alpha \rightarrow \gamma\beta$
 - Transitive Rule : If $\alpha \rightarrow \beta$ and $\beta \rightarrow \gamma$ then $\alpha \rightarrow \gamma$

$F^+ = F$

repeat

for each functional dependency f in F^+

 apply reflexivity and augmentation rules on f

 add the resulting functional dependencies to F^+

for each pair of functional dependencies f_1 and f_2 in F^+

if f_1 and f_2 can be combined using transitivity

 add the resulting functional dependency to F^+

until F^+ does not change any further

- Correctness of the algorithm ?
- Running time of the algorithm ?

Recap : Closure of Attribute set

- An Attribute B is functionally dependent on α if $\alpha \rightarrow B$
- To check if α is a superkey, we should compute the set of all attributes functionally dependant on α
- One way to do this is to compute F^+ , take all FDs with α as left-hand side and right-hand side will be the union of the right-hand side of all such FDs
 - This is expensive

Recap : Closure of Attribute set

- There is an efficient algorithm

```
result :=  $\alpha$ ;  
repeat  
  for each functional dependency  $\beta \rightarrow \gamma$  in  $F$  do  
    begin  
      if  $\beta \subseteq result$  then  $result := result \cup \gamma$ ;  
    end  
until ( $result$  does not change)
```

$A \rightarrow B$

$A \rightarrow C$

$CG \rightarrow H$

$CG \rightarrow I$

$B \rightarrow H$

- Over schema $r(A,B,C,G,H,I)$ Compute $(AG)^+$
 - $result = \{A,G\}$
 - $result = \{A,G,B\}$
 - $result = \{A,G,B,C\}$
 - $result = \{A,G,B,C,H\}$
 - $result = \{A,G,B,C,H,I\}$

Recap : Closure of Attribute set - Uses

- To test if α is a superkey
- Check if a functional dependency $\alpha \rightarrow \beta$ holds
- Alternative way to compute F^+

Recap : Canonical Cover

- Suppose database has functional dependency set F :
- Whenever a database is updated, the system should ensure all functional dependencies are satisfied
- If we have a minimal set of functional dependencies F^* that imply all the functional dependencies of F then the check can be made faster

Recap : Canonical Cover

- An attribute of a functional dependency is extraneous if we can remove it without changing the closure set of the functional dependencies.
- Consider a set of functional dependencies F and let $\alpha \rightarrow \beta$ be in F
 - Attribute A in α is extraneous if F logically implies $(F - \{\alpha \rightarrow \beta\}) \cup \{(\alpha - A) \rightarrow \beta\}$
 - Attribute A in β is extraneous if the set $(F - \{\alpha \rightarrow \beta\}) \cup \{\alpha \rightarrow (\beta - A)\}$ logically implies F

Recap : Canonical Cover

- An attribute of a functional dependency is extraneous if we can remove it without changing the closure set of the functional dependencies.
- Consider a set of functional dependencies F and let $\alpha \rightarrow \beta$ be in F
 - Attribute A in α is extraneous if F logically implies $(F - \{\alpha \rightarrow \beta\}) \cup \{(\alpha - A) \rightarrow \beta\}$
 - Attribute A in β is extraneous if the set $(F - \{\alpha \rightarrow \beta\}) \cup \{\alpha \rightarrow (\beta - A)\}$ logically implies F
- Example:
 - In $F = \{AB \rightarrow C, A \rightarrow C\}$, B is extraneous in $AB \rightarrow C$
 - In $F = \{AB \rightarrow CD, AB \rightarrow C\}$ C is extraneous in $AB \rightarrow CD$
- Note: In both cases the other direction of the implication trivially holds always

Recap : Canonical Cover

Checking for extraneous attributes

- Consider a set of functional dependencies F and let $\alpha \rightarrow \beta$ be in F
 - Attribute A in α is extraneous if F logically implies $(F - \{\alpha \rightarrow \beta\}) \cup \{(\alpha - A) \rightarrow \beta\}$
 - Attribute A in β is extraneous if the set $(F - \{\alpha \rightarrow \beta\}) \cup \{\alpha \rightarrow (\beta - A)\}$ logically implies F
- For a relational schema R and functional dependencies F and $\alpha \rightarrow \beta$ in F ,
Let A be an attribute in $\alpha \rightarrow \beta$
 - If A is in β : To check if A is extraneous, consider the set of functional dependencies:
 $F^* = (F - \{\alpha \rightarrow \beta\}) \cup \{\alpha \rightarrow (\beta - A)\}$ and check if $\alpha \rightarrow \beta$ can be inferred from F^* (By computing α^+ under F^*)
 - If A is in α : To check if A is extraneous, consider $\gamma = \alpha - \{A\}$ and check if $\gamma \rightarrow \beta$ can be inferred under F (By computing γ^+)

Recap : Canonical Cover

- **Canonical Cover** of a set of functional dependencies F is given by F_c such that F logically implies F_c and vice-versa, such that the following conditions hold:
 - No functional dependencies in F_c contains extraneous attributes
 - Each left side of a functional dependency in F_c is unique(There are no two dependencies $\alpha_1 \rightarrow \beta_1$ and $\alpha_2 \rightarrow \beta_2$ in F_c such that $\alpha_1 = \alpha_2$)

$F_c = F$

repeat

Use the union rule to replace any dependencies in F_c of the form

$\alpha_1 \rightarrow \beta_1$ and $\alpha_1 \rightarrow \beta_2$ with $\alpha_1 \rightarrow \beta_1 \beta_2$.

Find a functional dependency $\alpha \rightarrow \beta$ in F_c with an extraneous attribute either in α or in β .

/* Note: the test for extraneous attributes is done using F_c , not F */

If an extraneous attribute is found, delete it from $\alpha \rightarrow \beta$ in F_c .

until (F_c does not change)

Recap : Canonical Cover

Example

- Let $r(A, B, C)$ be relational schema with the following functional dependencies:
$$\begin{array}{l} A \rightarrow BC \\ B \rightarrow C \\ A \rightarrow B \\ AB \rightarrow C \end{array}$$
- Step 0: $F_c = \{ A \rightarrow BC, B \rightarrow C, A \rightarrow B, AB \rightarrow C \}$
- Step 1 : Combine 1st and 3rd FD and obtain $F_c = \{ A \rightarrow BC, B \rightarrow C, AB \rightarrow C \}$
- Step 2: A is extraneous in $AB \rightarrow C$
 - because F_c implies $\{ B \rightarrow C, A \rightarrow BC \} \cup \{ B \rightarrow C \}$
 - $F_c = \{ A \rightarrow BC, B \rightarrow C \}$
- Step 3 : C is extraneous in $A \rightarrow BC$
 - Because $\{ A \rightarrow B, B \rightarrow C \}$ logically implies $A \rightarrow BC$
 - $F_c = \{ A \rightarrow B, B \rightarrow C \}$

Recap : Canonical Cover

- **Canonical Cover** of a set of functional dependencies F is given by F_c such that F logically implies F_c and vice-versa, such that the following conditions hold:
 - No functional dependencies in F_c contains extraneous attributes
 - Each left side of a functional dependency in F_c is unique
(There are no two dependencies $\alpha_1 \rightarrow \beta_1$ and $\alpha_2 \rightarrow \beta_2$ in F_c such that $\alpha_1 = \alpha_2$)
- **F and F_c** have same closures.
- Testing if F is satisfied is equivalent to testing if F_c is satisfied
- It is cheaper to test F_c than F itself

Dependency Preservation

- Let F be a set of functional dependencies on R and let $R_1 R_2 \dots R_n$ be a decomposition of R
- Restriction of F to R_i is the set F_i is the set of all functional dependencies that includes only the attributes of R_i
- Example: If $r(A,B,C)$ is a schema with $F = \{ A \rightarrow B, B \rightarrow C, A \rightarrow C \}$ and the decomposition is $R_1 = (A,B)$ and $R_2 = (A,C)$
Then $F_1 = \{ A \rightarrow B \}$ and $F_2 = \{ A \rightarrow C \}$
- Decomposition of R to $R_1 R_2 \dots R_n$ is said to be dependency preserving if:

$$F^+ = (F_1 \cup F_2 \cup \dots \cup F_m)^+$$

```
compute  $F^+$ ;  
for each schema  $R_i$  in  $D$  do  
  begin  
     $F_i :=$  the restriction of  $F^+$  to  $R_i$ ;  
  end  
 $F' := \emptyset$   
for each restriction  $F_i$  do  
  begin  
     $F' = F' \cup F_i$   
  end  
compute  $F'^+$ ;  
if ( $F'^+ = F^+$ ) then return (true)  
  else return (false);
```

Decomposition Algorithms : BCNF

- A relation schema R is in BCNF with respect to a set F of functional dependencies if, for all functional dependencies in F^+ of the form $\alpha \rightarrow \beta$, where $\alpha \subseteq R$ and $\beta \subseteq R$: at least one of the following holds:
 - $\alpha \rightarrow \beta$ is a trivial functional dependency (that is, $\beta \subseteq \alpha$)
 - α is a superkey for schema R
- Testing for BCNF
 - Goal is to see if there exists a non-trivial $\alpha \rightarrow \beta$ in F^+ where α is not a superkey
 - For every non-trivial $\alpha \rightarrow \beta$ in F^+ Compute α^+ and see if it contains all the attributes of R
 - There are Other clever ways to do it that avoids computing F^+

Decomposition Algorithms : BCNF

- A relation schema R is in BCNF with respect to a set F of functional dependencies if, for all functional dependencies in F^+ of the form $\alpha \rightarrow \beta$, where $\alpha \subseteq R$ and $\beta \subseteq R$: at least one of the following holds:
 - $\alpha \rightarrow \beta$ is a trivial functional dependency (that is, $\beta \subseteq \alpha$)
 - α is a superkey for schema R
- BCNF Decomposition Algorithm

```
result := {R};  
done := false;  
compute  $F^+$ ;  
while (not done) do  
    if (there is a schema  $R_i$  in result that is not in BCNF)  
        then begin  
            let  $\alpha \rightarrow \beta$  be a nontrivial functional dependency that holds  
            on  $R_i$  such that  $\alpha \rightarrow R_i$  is not in  $F^+$ , and  $\alpha \cap \beta = \emptyset$ ;  
            result := (result -  $R_i$ )  $\cup$  ( $R_i - \beta$ )  $\cup$  ( $\alpha, \beta$ );  
        end  
    else done := true;
```

Decomposition Algorithms : BCNF

- Example:
class (course_id, title, dept_name, credits, sec_id, semester, year, building, room_number, capacity, time_slot_id)
- course_id → title, dept_name, credits
- building, room_number → capacity
- course_id, sec_id, semester, year → building, room_number, time_slot_id
- Candidate Key = { course_id, sec_id, semester, year }
- course_id → title, dept_name, credits (is non-trivial FD and course_id is not superkey)
 - course (course_id , title, dept_name, credits)
 - class1 (course_id, sec_id, semester, year, building, room_number, capacity, time_slot_id)
- building, room_number → capacity (is non-trivial FD and course_id is not superkey)
 - course (course_id , title, dept_name, credits)
 - class11 (building, Room_number, capacity)
 - class12 (course_id, sec_id, year, time_slot_id)
- All 3 are in BCNF (This is lossless and dependency preserving decomposition)
- Algorithm ensures losslessness but dependency preservation might be violated

Decomposition Algorithms : 3NF

- A relation schema $r(R)$ is in Third Normal Form (3NF) with respect to a set F of functional dependencies if, for all functional dependencies in F^+ of the form $\alpha \rightarrow \beta$, where $\alpha \subseteq R$ and $\beta \subseteq R$: at least one of the following holds:
 - $\alpha \rightarrow \beta$ is a trivial functional dependency (that is, $\beta \subseteq \alpha$).
 - α is a superkey for schema R .
 - Each attribute A in $\beta - \alpha$ is contained in some candidate key of R

```
let  $F_c$  be a canonical cover for  $F$ ;  
 $i := 0$ ;  
for each functional dependency  $\alpha \rightarrow \beta$  in  $F_c$   
   $i := i + 1$ ;  
   $R_i := \alpha \beta$ ;  
if none of the schemas  $R_j, j = 1, 2, \dots, i$  contains a candidate key for  $R$   
  then  
     $i := i + 1$ ;  
     $R_i :=$  any candidate key for  $R$ ;  
  /* Optionally, remove redundant relations */  
repeat  
  if any schema  $R_j$  is contained in another schema  $R_k$   
  then  
    /* Delete  $R_j$  */  
     $R_j := R_i$ ;  
     $i := i - 1$ ;  
until no more  $R_j$ s can be deleted  
return  $(R_1, R_2, \dots, R_i)$ 
```

Decomposition Algorithms : 3NF

- Example:

dept_advisors (s_ID, i_ID, dept_name)

- $i_ID \rightarrow dept_name$

- $s_ID, dept_name \rightarrow i_ID$

- $F = F_c$ (no extraneous attributes)

- $R_1(i_ID, dept_name)$

- $R_2(s_ID, dept_name, i_ID)$

- Delete R_1 and retain only R_2

- Algorithm results in Lossless and Dependency preserving 3NF decomposition

let F_c be a canonical cover for F ;

$i := 0$;

for each functional dependency $\alpha \rightarrow \beta$ in F_c

$i := i + 1$;

$R_i := \alpha \beta$;

if none of the schemas R_j , $j = 1, 2, \dots, i$ contains a candidate key for R

then

$i := i + 1$;

$R_i :=$ any candidate key for R ;

 /* Optionally, remove redundant relations */

repeat

if any schema R_j is contained in another schema R_k

then

 /* Delete R_j */

$R_j := R_i$;

$i := i - 1$;

until no more R_j s can be deleted

return (R_1, R_2, \dots, R_i)

Comparing BCNF and 3NF

- It is **always possible to obtain 3NF** without sacrificing dependency preservation and losslessness
 - But might lead to having many nulls and repetition of information
- Goal:
 - BCNF + Lossless decomposition + Dependency Preservation
 - All 3 are not possible simultaneously
 - Choose BCNF or Dependency Preservation + 3NF
- SQL does not allow specification of FDs (except primary key and super keys using unique)
 - Can be done using assertions

Multivalued dependencies

- In some cases BCNF does not reduce repetitions
- Suppose one instructor can be associated with multiple departments. Moreover Instructors can have several addresses (office, current address, permanent address..)
 - inst (ID, dept_name, name, street, city)
 - FD : ID \rightarrow name (but ID is not the key for inst)
 - r_1 (ID, dept_name, street, city) and r_2 (ID, name)
- Decomposition is in BCNF but there still can be repetition in r_1
 - r_{11} (ID, dept_name) and r_{12} (ID, street, city)
- How to identify this in the absence of additional constraints ?

Multivalued dependencies

- Let $r(R)$ be a relational schema and let $\alpha \subseteq R$ and $\beta \subseteq R$.
The multivalued dependency $\alpha \twoheadrightarrow \beta$ holds if for every legal instance r of $r(R)$ and tuples t_1 and t_2 of r such that $t_1[\alpha] = t_2[\alpha]$ then there exists t_3 and t_4 such that

- $t_1[\alpha] = t_2[\alpha] = t_3[\alpha] = t_4[\alpha]$
- $t_3[\beta] = t_1[\beta]$
- $t_3[R - \beta] = t_2[R - \beta]$
- $t_4[\beta] = t_2[\beta]$
- $t_4[R - \beta] = t_1[R - \beta]$

	α	β	$R - \alpha - \beta$
t_1	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$a_{j+1} \dots a_n$
t_2	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$b_{j+1} \dots b_n$
t_3	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$b_{j+1} \dots b_n$
t_4	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$a_{j+1} \dots a_n$

- Multivalued dependencies are part of more general **tuple generating dependencies**
- Functional dependencies are part of more general **equality generating dependencies**

Multivalued dependencies : Closure

- If D is a set of FDs + MVDs then let D^+ denote the closure of D that contains all the dependencies that are logically implied by D
- Computed using axioms. Sample axioms:
 - If $\alpha \twoheadrightarrow \beta$ then $\alpha \twoheadrightarrow \beta$
 - If $\alpha \twoheadrightarrow \beta$ then $\alpha \twoheadrightarrow R - \alpha - \beta$

Multivalued dependencies : Fourth Normal Form

- MVD $\alpha \twoheadrightarrow \beta$ over $r(R)$ is trivial if $\beta \subseteq \alpha$ or $\alpha \cup \beta = R$
- A relational schema $r(R)$ is in Fourth Normal Form (4NF) with respect to a set D of FD + MVD if for all MVDs $\alpha \twoheadrightarrow \beta$ in D^+ where $\alpha \subseteq R$ and $\beta \subseteq R$ one of the following holds:
 - $\alpha \twoheadrightarrow \beta$ is a trivial MVD
 - α is a superkey of R
- A database design is in 4NF if every relational schema is in 4NF
- Note that 4NF implies BCNF

Multivalued dependencies : 4NF decomposition

- Preserves Losslessness
- But does not preserve dependencies

```
result := {R};  
done := false;  
compute  $D^+$ ; Given schema  $R_i$ , let  $D_i$  denote the restriction of  $D^+$  to  $R_i$   
while (not done) do  
    if (there is a schema  $R_i$  in result that is not in 4NF w.r.t.  $D_i$ )  
        then begin  
            let  $\alpha \twoheadrightarrow \beta$  be a nontrivial multivalued dependency that holds  
            on  $R_i$  such that  $\alpha \rightarrow R_i$  is not in  $D_i$ , and  $\alpha \cap \beta = \emptyset$ ;  
            result := (result -  $R_i$ )  $\cup$  ( $R_i - \beta$ )  $\cup$  ( $\alpha, \beta$ );  
        end  
    else done := true;
```

Higher Normal Forms

- Join dependencies generalize multivalued dependencies
 - Leads to Project-Join Normal Form (5NF)
- Domain Key Normal Form
- Disadvantages of higher normal forms:
 - Hard to reason
 - No sound and complete axiom set

Database Design and Normalization

- Generated while converting an E-R diagram into a relational schema
 - **Good E-R diagram** does not need too much normalization when translated to relational schemas
 - **Relationship sets involving more than two entities** can result in non-normalized schema
 - **Universal-relation approach** : Design one big relational schema that contains everything and then normalize
- Naming Relationship and Attributes
 - **Unique Role assumption** : Every name of attribute has a unique meaning
 - Convention to designate primary key attributes at the beginning
 - **Prefix helps in disambiguation** : student_ID, instructor_ID

Database Design and Normalization

- Denormalization for Performance
 - Occasionally designers choose to retain redundancy
 - Improves performance
(Normalized databases need join operations to evaluate queries)
 - Alternative is to use normalized database and store denormalized database as materialized views

Database Design and Normalization

- Some aspects not considered in normalization process
- Example:
 - Store number of instructors in each department in various years
`(dept_name, year, num_inst)`
with `dept_name, year → num_inst` as FD
 - Alternative is to have one table for each year
`inst_count_2020 (dept_name num_inst)` `inst_count_2021 (dept_name, num_inst)`
With `dept_name → num_inst` as FD in each table
 - Another Alternative
`dept_year (dept_name, num_inst_2020, num_inst_2021)`
With `dept_name → num_inst_2020, num_inst_2021`
Close to spreadsheet representation
- Last 2 alternatives are bad design choices though they are in BCNF

Modelling Temporal Data

- Temporal Data is associated with a time interval where the data is valid (called snapshot)
- Functional dependency should take temporal aspect into account
- Example : (ID, street, city, from, to)
- Temporal functional dependency : $X \xrightarrow{T} Y$
 - Makes the design and functional dependence analysis complicated
 - First design without temporal data (by considering a snapshot) then decide what temporal constraints are needed
- What should be the end time if it is currently ongoing?
 - Null ? (Problem with primary key)
 - Something far in future
- Relationship between time dependent and time independent data
 - Advisor relationship and department budget (though depends on current budget)
 - Solution: Not to add time in main relation, but maintain history separately

Reference:

Database System Concepts by Silberschatz, Korth and Sudarshan
(6th edition)
Chapter 8